

## Reminders

### Upcoming due dates

TODAY: Teaching evaluations for extra credit!

Wed Mar 18 Final project\*, video\*, team eval survey, post-course survey

# Your DS future

## Data Science in Practice

Jason G. Fleischer, PhD

Dept. of Cognitive Science

UC San Diego

# Courses in DS and ML at UCSD

**Did you know about the reciprocal agreement between DS and COGS??**

- DS [My list of '20-21 ML \(and ML adjacent\) courses](#)
- CSE
- CS
- ECE
- COGS
- But also many other departments like ECON, MATH, LING, BENG, etc

# Some job titles and what they do

- Analytics or statistician: data handling, analysis
- Data scientist: programming, data handling, analysis
- Data engineer: programming, databases, management
- Data architect: programming, databases, design
- Data manager: databases, design, management
- \*OPs (eg, devOPs, dataOPs, full stack): programming, tool development, management concentrating on end to end process
- ML Engineer: programming, tool development, management of infrastructure
- ML researcher: programming, algorithm design and testing



| New Roles and Job Responsibilities in 2024 from Burtchworks DS report |                                                                                                                                                                                                                           |
|-----------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                                       |                                                                                                                                                                                                                           |
| <b>AI Ethics Officer:</b>                                             | Develops, implements, and oversees ethical guidelines and practices for AI and machine learning systems. Ensures that AI applications are transparent, fair, and free from biases.                                        |
| <b>Machine Learning Operations (MLOps) Engineer:</b>                  | Focuses on deploying, monitoring, and maintaining machine learning models in production. Collaborates with data scientists and IT teams to automate and streamline ML workflows.                                          |
| <b>Data Science Product Manager:</b>                                  | Manages the development and launch of data-driven products. Works with data scientists, engineers, and stakeholders to ensure product alignment with business goals and user needs.                                       |
| <b>AI Research Scientist:</b>                                         | Conducts cutting-edge research in artificial intelligence, focusing on developing new algorithms, models, and techniques. Publishes findings in scientific journals and presents at conferences.                          |
| <b>Deep Learning Engineer:</b>                                        | Specializes in building and optimizing deep neural networks for tasks such as image recognition, natural language processing, and speech recognition. Works with large datasets and high-performance computing resources. |
| <b>Natural Language Processing (NLP) Engineer:</b>                    | Develops and implements algorithms for understanding and generating human language. Works on applications like chatbots, sentiment analysis, and language translation.                                                    |
| <b>AI Solution Architect:</b>                                         | Designs and implements AI solutions to solve complex business problems. Works closely with stakeholders to understand requirements and develop scalable AI architectures.                                                 |
| <b>Data Privacy Officer:</b>                                          | Ensures that data handling practices comply with privacy laws and regulations. Works on protecting sensitive information and developing data governance policies.                                                         |
| <b>Computer Vision Engineer:</b>                                      | Focuses on developing algorithms and systems for analyzing and interpreting visual data from the real world. Works on applications like autonomous vehicles, facial recognition, and medical imaging.                     |
| <b>AI Educator/Trainer:</b>                                           | Provides training and education on AI and machine learning topics. Develops curriculum and delivers courses, workshops, and seminars to students, professionals, and organizations.                                       |
| <b>Robotics Engineer:</b>                                             | Designs, builds, and programs robots and robotic systems. Works on integrating AI and machine learning to enhance robotic capabilities and autonomy.                                                                      |
| <b>Data Storyteller:</b>                                              | Translates complex data analysis into compelling narratives and visualizations. Communicates insights to stakeholders in a clear and impactful manner.                                                                    |
| <b>Cognitive AI Developer:</b>                                        | Develops AI systems that mimic human thought processes. Works on applications such as decision support systems, cognitive automation, and AI-based problem-solving tools.                                                 |
| <b>Federated Learning Engineer:</b>                                   | Specializes in federated learning, where machine learning models are trained across multiple decentralized devices or servers. Focuses on privacy-preserving and collaborative AI.                                        |
| <b>AI Quality Assurance (QA) Specialist:</b>                          | Ensures the quality and performance of AI models and systems. Conducts testing, validation, and evaluation to identify and resolve issues before deployment.                                                              |
| <b>Autonomous Systems Engineer:</b>                                   | Develops and maintains autonomous systems, such as self-driving cars or drones. Integrates AI algorithms to enable autonomous decision-making and operation.                                                              |
| <b>AI Personalization Specialist:</b>                                 | Focuses on creating personalized user experiences using AI. Works on recommendation systems, adaptive learning platforms, and personalized marketing.                                                                     |
| <b>AI-driven Healthcare Specialist:</b>                               | Develops AI applications for the healthcare industry. Works on predictive analytics, medical imaging analysis, and personalized treatment plans.                                                                          |
| <b>AI Regulation and Compliance Specialist:</b>                       | Ensures that AI systems comply with regulatory requirements and standards. Works on developing frameworks for legal and ethical AI use.                                                                                   |
| <b>Voice and Speech Recognition Engineer:</b>                         | Develops algorithms and systems for recognizing and synthesizing human speech. Works on applications like virtual assistants, transcription services, and voice biometrics.                                               |



# What should you learn next?

- Data science?
  - Statistics + Modeling
  - SQL and/or Data viz packages
  - Cloud!
  - Become fantastic at one of Python/R and learn many packages. Be capable at the other language
- AI?
  - Statistics, Vector calculus, ML theory
  - Cloud!
  - Foundation models
  - Become usefully good at several frameworks (PyTorch, Go, Tensorflow, etc)
  - Read arXiv and NeurIPS/ICLR papers
- Both: learn to properly use AI tools to help you :^/

# How do you use AI

- Windsurf and Cursor... take your repo and modify it .. local
- Whisper: spoken vibe coding
- Using whatever LLM for research
- Use Ai to help figure out a possible solution, then look up docs, think about how to use the function to get the result
-



- Mostly debugging... explain an error
- Onboard team members... specialize a task for one team member
- Researching.. brief summary of a concept
- Documentation
- Refactoring your code
- Organizing report or code
- Brainstorming partner
- Practice problems
-

# How to properly use AI tools

## As a programmer

- Beginner
  - Define a narrow task like “write a function to individually z-score only the numeric columns of a Pandas dataframe. The operation should not modify the data frame in place, it should generate a copy of the original data frame and return that copy “
  - Ask AI to explain errors to you
  - Ask it to suggest better approaches to a problem
  - Ask it to write a unit test, or in chat ask it to back and forth with you on what kinds of tests make sense
- Agentic approaches
  - Make a detailed English planning document in a .md file
  - Set the agents to work on different parts of it, including allowing them to modify and expand the plan
  - e.g. <https://kiro.dev/>

# How to properly use AI tools

## As a data scientist

- Define a narrow task like “write a function to individually z-score only the numeric columns of a Pandas dataframe. The operation should not modify the data frame in place, it should generate a copy of the original data frame and return that copy”
- Ask it to help with writing API calls to accomplish something
- Ask it to help with data cleaning or generating summary reports
- You need to supply the deep thinking and analysis... its a tool to get it done more efficiently!



# **2025 Burtchworks Datascience and AI Professionals Salary Report**

# ROLE CLASSIFICATION & SAMPLE

**Burtch Works** differentiates these groups by their **primary mission and deliverables**, not just data modality.


**Data Scientists** focus on driving decisions with statistical analysis and classical ML. They own experimentation and causal inference, forecasting, segmentation, and predictive modeling on business/operational data; their outputs are analyses, validated experiments, and models embedded in analytics/BI or batch/stream scoring.

**AI Professionals** design, build, and operate **ML/LLM-powered systems** that interact with users and workflows. They work across un/semi-structured and multimodal data; fine-tune and distill models; implement retrieval/grounding; orchestrate tools/agents; run RLHF & safety evaluations; and manage MLOps (serving, monitoring, cost/latency, governance).

Roles increasingly overlap; we classify each professional by what they primarily do day-to-day in their current job.

**2025 Sample**  
724 Data Scientists and 162 AI Professionals (total **886**), based on a strict 12-month window and the same segmentation used in prior years.

Burtch Works has tracked Data Science and AI compensation for over a decade. The 2025 cycle (May 2024 – April 2025; **724 Data Scientists** and **162 AI Professionals**) shows a market that’s moved from experimentation to **operational AI** — and pay is adjusting to where the scarcest, most business-critical skills now sit.







Where we are  
**TODAY**

*AI & Data Science Hiring Landscape*

## FROM HYPE TO HANDS-ON

**Gen-AI moves from pilots to production**  
Hiring shifts from “prompt roles” to **Applied LLM Engineers** who can integrate, monitor, and cost-optimize real workflows.

**RLHF & alignment stay supply-constrained (+34% reqs vs +7% supply)**  
Scarcity drives premiums and pushes teams toward train-to-hire pipelines and selective offshore/nearshore augmentation.

**AI Innovation Labs**  
Labs are evolving into **productization engines** — owning reference architectures, governance, and sprint pods that turn POCs into shipped, ROI-tracked features.

**Toughest new-grad market in 10 years**  
Fewer research-track seats; entry roles skew to AI Operations or automation, making internships and portfolios the decisive hiring requirement.

**IC-led hiring dominates**  
Firms favor high-leverage ICs who can ship PoCs rapidly.

**Increased demand (but with more selective hiring)**  
AI-related job postings in the U.S. more than doubled from 66,000 to nearly 140,000 between January and April 2025. Despite the surge, economic uncertainty is prompting employers to be more selective.

**Broader business value and cross-functional demand**  
AI and Data Science hiring is expanding beyond engineering into growth, marketing, sales, and customer success functions.

*Key Takeaways*

## Q3/Q4 HIRING SURVEY

**75%**

of new hires are for **Data & Analytics** roles. Market Research, Engineering, and Cyber Security also reflect current enterprise priorities.

**60%**

of companies are **holding steady** with their current teams, indicating a conservative hiring posture, while nearly **20%** are downsizing.

**68%**

of new hires are for **full-time roles**, showing a commitment to long-term workforce development.

**60%**

of new hires are for **backfills**, suggesting internal attrition or restructuring.

**Small-scale hiring**

of 1-5 roles is the norm, with very few large hiring initiatives.

**Technical skills**

are the top priority, followed by critical thinking, both outranking cultural fit and industry familiarity.



Hiring Demand  
**OVERVIEW**

*AI Matures, Delivers Value*

## AI MOVES FROM EXPERIMENTS TO SCALABLE PLATFORMS

**From pilots to platforms**  
Enterprises consolidate tools and standardize on an AI platform stack, shifting budgets from POCs to **repeatable, ROI-tracked features**.

**Cost-to-value discipline**  
Token, context, and inference costs are now product KPIs. Teams adopt **distillation, caching, PEFT**, and **small/efficient models** to hit latency and unit-economics targets.

**Role evolution**  
“Prompt engineer” titles fade into **Applied LLM Engineer, Retrieval Engineer**, and **AI Platform Engineer**. Demand rises for **RLHF/Alignment Ops** and **AI Evaluation/Safety Specialists**.

**Multimodal & agents mature**  
Voice/vision inputs and **tool-using agents** leave labs for service, sales, and ops — measured on containment, CSAT, and cost per resolution.

**Hybrid buy-build wins**  
Teams combine vendor services with **open-weights** and in-house fine-tuning for control, privacy, and portability.

**Top use cases**  
**Knowledge search/RAG, copilots for sales & service, analytics summarization, and workflow automation.** Low-ROI “general chatbots” decline.

**Big capital flow & infrastructure buildouts**  
Massive investments in data center infrastructure and workforce development. Private equity is fueling hyperscalers and co-locations to meet rising demand for AI compute as global data center spending could hit \$3 trillion in 2029.

**Geographical & regulatory shifts**  
The landmark “TAKE IT DOWN Act” criminalizes the publication of non-consensual intimate imagery (weather authentic or AI-generated). The One Big Beautiful Bill Act (BBB) allows companies to immediately deduct 100% of domestic R&D expenditures, including salaries, equipment, and related costs — a significant departure from the previous multi-year amortization requirement.



# TEAM COMPOSITION & TALENT STRATEGY

## Gen-AI Implementation (Applied LLM in production)

### Role families & representative titles

- Applied LLM Engineer, Retrieval Engineer, Agent/Orchestration Engineer
- AI Solutions Architect (Apps), AI Product Engineer

### What they do

- Ship user-facing copilots/agents and RAG apps; design tool-use/function-calling, memory, and guardrails.
- Optimize chunking, embeddings, rerankers, and grounding data quality.

### Must-have skills

- Python/TypeScript; LangChain/LlamaIndex or equivalent; vector DBs; PEFT/LoRA; distillation.
- Evaluation & observability (quality, latency, cost per resolution); prompt hardening.

### KPIs

- Containment %, CSAT, tasks/hour, \$ per task/token, latency p95, A/B win-rate vs baselines.

### 2025 market note

- Hiring shifts from “prompt engineer” titles to Applied LLM / Retrieval / Agent roles; entry-level comp up, mid-level AI bases flat/down.

## Data Science & Analytics

### Role families & representative titles

- Product/Decision Data Scientist, Causal Inference Scientist, Forecasting/Time-Series Scientist
- Marketing/Revenue Analytics, Fraud & Risk, Experimentation Lead

### What they do

- Drive decisions and measurable lift: experimentation, causal models, forecasting, pricing, segmentation.
- Partner with product/ops to productionize models into BI, batch/stream scoring, or features.

### Must-have skills

- A/B and quasi-experimental design, uplift models, time-series, feature engineering, SQL/Python.
- Communicating impact; cost/benefit and instrumentation.

### KPIs

- Incremental revenue/savings, test velocity, forecast error, adoption of DS outputs.

### 2025 market note

- DS leadership (MG-2/MG-3) re-priced up as orgs operationalize analytics and govern AI programs.

# TEAM COMPOSITION & TALENT STRATEGY

## AI-Lab Build-Out

### Role families & representative titles

- Head of AI / Lab Director (Director/VP), AI Product Lead, Technical Program Manager (AI)
- AI Solutions Architect (LLM/RAG), Governance & Risk Lead (MRM), Data Sourcing Lead
- Evaluation/Safety Lead, Labeling/Ops Manager

### What they do (2025)

- Stand up a reference architecture (data prep → retrieval → serving → eval) and an operating model (charter, intake, governance).
- Prioritize use cases, build the roadmap, staff pods, and own vendor/open-weights strategy.
- Institutionalize evaluation, safety, privacy, and spend controls from day one.

### Must-have skills

- Translating business goals into AI roadmaps; RAG/agents patterns; platform thinking.
- Model risk management (MRM), policy, data lineage/consent; auditability.
- Org design & hiring: sequencing IC-heavy pods before adding management layers.

### KPIs

- Time-to-first-POC / time-to-prod, adoption, governance pass-rates, cost per task / latency SLOs.

### 2025 market note

- CoE leaders with ops + governance chops command a premium; many orgs delay VP hires but will pay for a Director who can ship.

## RLHF & Alignment

### Role families & representative titles

- RLHF Trainer / Rater (Onshore/Nearshore), RLHF Ops Manager, Preference Model Engineer
- Safety/Policy Engineer, Red-Team Lead, Evaluation Engineer

### What they do

- Design guidelines, collect preference data, train reward/preference models, and run DPO/ORPO pipelines.
- Operationalize safety: policy taxonomies, automated eval harnesses, and incident response.

### Must-have skills

- Sampling/label quality control; reward modeling; safety eval (toxicity, bias, leakage).
- Tooling: label platforms, prompt/eval frameworks, LLM Guardrails, synthetic data generation.

### KPIs

- Aligned-response rate, policy-violation rate, reward model accuracy, data throughput/cost.

### 2025 market note

- Supply-constrained domestically; teams blend onshore leads with nearshore scale to meet SLAs.

# TEAM COMPOSITION & TALENT STRATEGY

## MLOps & Platform Engineering

### Role families & representative titles

- ML Platform Engineer, Model Reliability/SRE (AI), Data/Streaming Engineer
- Evaluation Platform Engineer, Model Governance Engineer

### What they do

- Provide the paved-road: CI/CD for ML/LLM, feature stores, model registry, inference gateways.
- Production observability (quality, drift, latency, safety) and cost controls; multi-cloud portability.

### Must-have skills

- Kubernetes, GPUs, Triton/Transformers inference, Ray; eval pipelines; policy & PII controls.
- Infra as code; telemetry & alerting; blue/green and canaries for models.

### KPIs

- Deployment frequency, lead time, change-fail rate, MTTR; drift incidents; infra \$ per request.

### 2025 market note

- “Buy the platform, build the fit” hybrid is standard — engineers who stitch vendor + open-weights are prized.

## Chief Data Officer (CDO)

### Role families & representative titles

- Chief Data Officer, Head of Data & Analytics, SVP Data Strategy, Chief Analytics Officer, Chief Data Analytics Officer, Chief AI Officer, Chief Strategy Officer

### What they do

- Set enterprise-wide data strategy aligned to business goals, including governance, analytics, automation of low value actives, new role design and AI adoption.
- Drive data monetization initiatives, advanced analytics programs, and AI/ML integration into products, operations and business process.

### Must-have skills

- Advanced analytics, AI/ML strategy, and business use case identification.
- Strong leadership and change management skills to shift organizational culture toward data fluency.

### KPIs

- Adoption rate of analytics/AI tools across the enterprise.
- ROI on data initiatives (cost savings, revenue growth, enterprise value creation).

### 2025 market note

- Role has evolved from governance-first to value-creation-focused, with increasing emphasis on AI enablement and strategy.
- High demand in regulated industries (financial services, healthcare, manufacturing) and PE-backed growth companies.



*Data Engineering Roles*

# DESIRED SKILL SETS: THE BASICS

In 2025, fundamentals of Python, SQL, Java, and Scala remain the foundation.

Cloud fluency in AWS, Azure, and GCP is now assumed, while distributed-framework expertise in Apache Spark, Kafka, and Databricks has become universal.

Familiarity with Pandas, NumPy, and dbt is table stakes, and employers increasingly expect Docker, Kubernetes, and workflow orchestration tools such as Airflow.

The market also emphasizes data governance, lineage, and observability, as organizations seek to ensure trust and compliance across expanding AI-powered data ecosystems.

# DESIRED SKILL SETS: THE NEXT LEVEL

Organizations now prioritize engineers who can productionize AI pipelines — bridging data engineering with MLOps and DataOps, a core theme of this year’s report.

Experience with TensorFlow, PyTorch, and LLM integration is valued not for model training itself, but for enabling *model-ready data systems*.


Engineers capable of deploying CI/CD pipelines for data, integrating AI-assisted data-quality monitoring, and managing feature stores and metadata catalogs are commanding premium compensation and accelerated advancement.



# EDUCATION: COMPARISON OF DEGREE LEVEL

- **~91% of AI professionals hold a graduate degree (Master's or PhD) in 2025 vs ~86% in 2024.**
- **Master's now dominates:**
  - AI: Master's ≈ 64% in 2025 (up ~+19 pp YoY), overtaking PhD as the largest credential.
    - Master's has the sufficient expertise for using GenAI
  - Data Science: Master's ≈ 64% in 2025 (up ~+7 pp YoY).
- **Bachelor's share keeps shrinking:**
  - AI: ≈ 8% in 2025 (down from ~14%).
  - Data Science: ≈ 16% in 2025 (down from ~24%).
- **Implication:** Entry-level AI jobs increasingly demand a graduate credential or equivalent experience (bootcamps, high-profile open-source contributions). For DS roles, a Bachelor's remains viable, but the credential ceiling is rising.
- **Implication:** Unless you are working at the big labs or FAANG, a Master's degree is enough though a PHD degree may still be considered.

Comparison of Degree Level: Data Science vs AI (2024 vs 2025)



EDUCATION:  
SALARIES BY JOB LEVEL FOR  
DATA SCIENCE PROFESSIONALS


| Job Level | Degree | 25th      | Mean      | 75th      |
|-----------|--------|-----------|-----------|-----------|
| IC-1      | BS     | \$72,500  | \$94,643  | \$110,000 |
|           | MS     | \$88,750  | \$110,027 | \$130,000 |
|           | PhD    | \$127,500 | \$131,818 | \$137,500 |
| IC-2      | BS     | \$115,000 | \$136,994 | \$150,000 |
|           | MS     | \$120,000 | \$138,421 | \$150,000 |
|           | PhD    | \$130,000 | \$148,146 | \$160,000 |
| IC-3      | BS     | \$130,000 | \$150,888 | \$170,000 |
|           | MS     | \$140,000 | \$159,073 | \$177,500 |
|           | PhD    | \$150,000 | \$163,021 | \$180,000 |
| MG-1      | BS     | \$130,000 | \$151,045 | \$173,500 |
|           | MS     | \$150,000 | \$162,496 | \$180,000 |
|           | PhD    | \$175,000 | \$183,333 | \$200,000 |
| MG-2      | BS     | \$163,775 | \$190,156 | \$200,000 |
|           | MS     | \$170,100 | \$194,030 | \$206,500 |
|           | PhD    | \$172,550 | \$189,704 | \$200,000 |
| MG-3      | BS     | \$242,750 | \$263,833 | \$300,000 |
|           | MS     | \$200,000 | \$239,895 | \$250,000 |
|           | PhD    | \$200,000 | \$252,871 | \$293,750 |

EDUCATION:  
SALARIES BY JOB LEVEL FOR  
AI PROFESSIONALS

| Job Level | Degree | 25th      | Mean      | 75th      |
|-----------|--------|-----------|-----------|-----------|
| IC-1      | BS     | \$146,250 | \$152,500 | \$158,750 |
|           | MS     | \$110,000 | \$130,063 | \$150,000 |
|           | PhD    | \$110,000 | \$120,000 | \$130,000 |
| IC-2      | BS     | \$110,000 | \$120,000 | \$130,000 |
|           | MS     | \$130,000 | \$149,736 | \$165,000 |
|           | PhD    | \$120,050 | \$145,743 | \$165,000 |
| IC-3      | BS     | \$140,000 | \$173,040 | \$185,000 |
|           | MS     | \$160,000 | \$165,623 | \$180,000 |
|           | PhD    | \$150,000 | \$175,891 | \$200,050 |
| MG-1      | MS     | \$142,500 | \$155,000 | \$167,500 |
|           | PhD    | \$147,500 | \$165,000 | \$182,500 |
| MG-2      | BS     | \$215,050 | \$223,367 | \$235,000 |
|           | MS     | \$177,525 | \$204,192 | \$212,500 |
|           | PhD    | \$170,000 | \$177,914 | \$180,050 |
| MG-3      | BS     | \$250,000 | \$250,000 | \$250,000 |
|           | MS     | \$200,000 | \$246,144 | \$275,000 |
|           | PhD    | \$225,000 | \$241,667 | \$272,500 |




# GENDER BREAKDOWN: DATA SCIENCE & AI




# CANDIDATE RESIDENCY STATUS

Companies appear more willing to hire non-US citizens (F1/OPT, lawful permanent residents, and H-1B visa holders) for AI roles, compared to Data Science positions.

Data Science Professionals Residency Status Breakdown



AI Professionals Residency Status Breakdown



# 2025 Burtchworks Research Professionals Salary Report



**Key Takeaways**

# 2025: DATA INTEGRATION, RESEARCH, AND ADVANCED TOOLS

Employers and professionals are navigating an evolving landscape shaped by the following factors:

**Integration of Data**

Market researchers are increasingly expected to merge primary research with secondary sources and big data, enabling a more comprehensive consumer view.

**Cross-disciplinary Collaboration**

Collaboration with data science, analytics, and technology teams is becoming the norm, particularly for projects requiring advanced insights and predictive capabilities.

**Adoption of Advanced Tools**

There is a growing emphasis on tools like Python, R, SQL, and machine learning models to enable deeper insights and foster automation within research processes.

**Increased Focus on ROI**

Companies are emphasizing measurable outcomes, requiring market researchers to align insights with clear business objectives and demonstrate return on investment (ROI) for research initiatives.

**Expansion of DIY Research Platforms**

Tools like Qualtrics and SurveyMonkey are enabling organizations to conduct quick-turn studies, although this trend emphasizes the need for professionals to guide proper design and analysis.



**Current Market Impacts**

# UPTICK IN HIRING AND HYBRID WORK MODELS

**Job Market Flux**

The market has shifted from an employer-driven hiring freeze in previous years to increased hiring activity in 2024. This uptick has generated optimism for sustained growth through 2025.

**Hybrid Work Models**

Hybrid working remains the most popular arrangement, with client-side roles typically requiring 2-3 days of in-person attendance. Supplier-side organizations continue to offer greater flexibility for remote work.

**Competitive Landscape**

While the candidate pool remains rich due to layoffs in 2023, companies are now actively seeking to retain top talent by streamlining hiring processes and offering attractive packages.





### Navigating the Market

# HOW TO NAVIGATE THE CURRENT HIRING MARKET



### Growing Importance of Data Integration

Market researchers are now expected to combine primary research with secondary and big data sources, creating a more holistic view of consumers.

### Cross-disciplinary Collaboration

Engagement with data science teams is becoming standard, so proficiency across these technical areas is a benefit for market research professionals.

### Emerging Tools and Platforms

New tools and platforms are emerging for automation, visualization, or AI-driven analysis. Market researchers should adopt more advanced tools and techniques such as Python, R, and SQL, and learn machine learning models to derive deeper insights from data.

Supplier side:  
Consultancy / X as a service


Client side:  
In-house for a company

| Supplier-Side |                             |
|---------------|-----------------------------|
| Level         | Typical Years of Experience |
| S-1           | 0-4 years                   |
| S-2           | 5-8 years                   |
| S-3           | 9-14 years                  |
| S-4           | 15+ years                   |


| Client-Side |                             |
|-------------|-----------------------------|
| Level       | Typical Years of Experience |
| C-1         | 0-4 years                   |
| C-2         | 5-8 years                   |
| C-3         | 9-14 years                  |
| C-4         | 15-24 years                 |
| C-5         | 25+ years                   |

# COMPENSATION CHANGES: 2023 AND 2024 BASELINE DATA

### Comparison of Supplier-Side Mean Base Salary: 2023 v. 2024



### Comparison of Client-Side Mean Base Salary: 2023 v. 2024



# Appen (aka Figure-Eight aka Crowdflower) State of AI report

## Methodology

The State of AI 2024 survey aimed to assess AI and data implementation across organizations. We partnered with The Harris Poll to conduct an online survey of 509 U.S. information technology decision-makers (ITDMs) from April 18 to May 9, 2024. Participants were ages 21 to 65 and worked at companies with 100+ employees.

### Industry

- Automotive & Transportation
- Finance & Financial Services
- Healthcare & Pharmaceuticals
- Internet & Electronics
- Retail & Consumer Durables (Including E-Commerce)
- Technology
- Telecommunications

### Role

- Business Leaders and Managers
- Data Scientists
- Data Engineers
- Developers

### Company sizes

- Small: 101-1,000 employees
- Medium: 1,001-10,000 employees
- Large: 10,001+ employees

*The research was conducted online in the United States by The Harris Poll on behalf of Appen among 509 information technology decision-makers. The survey was conducted April 18 - May 9, 2024. Raw data were not weighted and are therefore only representative of the individuals who completed the survey.*

*Respondents for this survey were selected from among those who have agreed to participate in our surveys. The sampling precision of Harris online polls is measured by using a Bayesian credible interval. For this study, the sample data is accurate to within  $\pm 4.3$  percentage points using a 95% confidence level. This credible interval will be wider among subsets of the surveyed population of interest.*

*All sample surveys and polls, whether or not they use probability sampling, are subject to other multiple sources of error which are most often not possible to quantify or estimate, including, but not limited to coverage error, error associated with nonresponse, error associated with question wording and response options and post-survey weighting and adjustments (not applicable in this case).*

### Small footnote

The 2023 study was conducted from February to March 2023 among 400 US and 100 European (UK, Ireland, and Germany) ITDMs. The 2022 study was conducted by The Harris Poll in June 2022 among 504 ITDMs; the 2021 study was conducted in March 2021 among 501 US ITDMs; and the 2020 study was conducted using a different method among 290 ITDMs.



# Introduction

As artificial intelligence (AI) adoption accelerates, the challenges surrounding data management, model reliability, and scalability are intensifying. In our 2024 State of AI report, we explore the critical factors driving AI development, including the rise of generative AI, enterprise headwinds, the importance of data quality, and the increasing demand for strategic partnerships to address the complexities of AI implementation. Our research, conducted in collaboration with The Harris Poll, surveyed over 500 U.S. information technology decision-makers (ITDMs) across a range of industries. The findings reveal five key takeaways shaping the future of AI:

1. **Generative AI Adoption is Increasing**  
- up 17 percentage points from 2023
2. **Enterprise Deployments & ROI are Down**  
- lower percentage of projects get deployed and show ROI
3. **Quality Data is Essential to AI Success**  
- 97% say data diversity, reduced bias, or scalability are crucial
4. **Data Challenges Persist**  
- 10 percentage point increase in data preparation bottlenecks
5. **Companies Seek Strategic Data Partners**  
- 93% want full AI Data Lifecycle expertise


These insights demonstrate that while AI offers transformative potential, its success depends on addressing ongoing data challenges and forming strong partnerships. At Appen, we're committed to supporting organizations through our comprehensive AI Data Platform (ADAP) and a global network of contributors, enabling businesses to navigate the complexities of AI with confidence.

| Use of Generative AI                                                                            | 2023 | 2024 |
|-------------------------------------------------------------------------------------------------|------|------|
| Use Generative AI to improve productivity and efficiency of internal business processes         | 49%  | 53%  |
| Use Generative AI to support internal IT operations                                             | 47%  | 49%  |
| Use Generative AI to aid in the research and development of new products                        | 37%  | 46%  |
| Use Generative AI to support other business functions (e.g., Marketing, Communications, Design) | 45%  | 43%  |
| Use Generative AI to support manufacturing/production operations                                | 36%  | 41%  |
| Use Generative AI based features into products or services that we sell                         | 40%  | 41%  |
| Use Generative AI to reduce business costs                                                      | 37%  | 41%  |

© 2024 State of AI Report **Appen**

Despite its rapid growth, managing bias and ensuring fairness in generative AI model training remains a key challenge. To tackle this, custom data collection has become the primary method for sourcing training data, highlighting its importance in creating reliable AI systems.

MOST COMMONLY USED DATA SET BY USE CASE




© 2024 State of AI Report **Appen**


# Data Challenges Persist

Sourcing quality data remains a significant challenge, becoming more difficult as AI use cases grow increasingly specialized. Data management has emerged as the leading obstacle in AI projects for ITDMs, with a reported 10 percentage point increase in bottlenecks related to sourcing, cleaning, and labeling data since 2023.

The lack of data availability has also risen by 7 percentage points during the same period. This increase may be attributed to the growing range of generative AI applications, which are driving new data requirements and aligned with the earlier observation of a rise in custom data collection for generative AI models.



Despite the importance of high-quality data annotation in AI development, average accuracy of data has dropped about 9 percentage points since 2021. This may be a result of the increasing complexity of AI systems and the corresponding data annotation requirements, which increasingly require specialized domain-specific knowledge and quality control across large datasets.



To overcome these challenges, organizations must shift their attention from short-term gains to long-term value creation in AI and focus on building a robust foundation of high-quality data.

# Quality Data is Essential to AI Success

**DATA DIVERSITY IS THE MOST VALUED FEATURE IN AI, FOLLOWED BY BIAS REDUCTION AND SCALABILITY.**

**80%** ACKNOWLEDGING THAT **HUMAN INSIGHT IS KEY**  
OF RESPONDENTS TO REFINING AI SYSTEMS

Our latest research reveals that a significant majority—86% of respondents—are retraining or updating their machine-learning models at least once every quarter. This underscores the rapid development across the industry and the close feedback loop between data and models in the AI training process.

Ninety percent of respondents confirmed that they rely on external data providers for their AI model training and/or annotation.

Custom-collected datasets, especially in text, images, video, audio, and point-of-interest data, have become the backbone of many leading AI applications. The effectiveness of these models depends on data that is accurate, diverse, and representative. Notably, 80% of respondents highlight the importance of human-in-the-loop machine learning, acknowledging that human insight is key to refining AI systems.

Data diversity is the most valued feature of AI data, closely followed by efforts to reduce bias, and ensuring data scalability. 7% of respondents agree these elements are vital for building AI models. Quality human-in-the-loop training and evaluation data play a crucial role in achieving these goals by providing the nuanced, diverse, and accurate data necessary to refine models, reduce bias, and ensure they perform reliably at scale.



**What else should I be thinking  
about in my career?**

# Soft skills

<https://www.manning.com/books/build-a-career-in-data-science>

<https://www.youtube.com/watch?v=pBisSm0Zr3A>

## BUILD A CAREER IN Data Science

Emily Robinson  
Jacqueline Nolis

 MANNING




# Problems with AI


New new

- Interpretability
- Alignment






# Paperclip AI Thought Experiment




I NEED  
LOTS OF  
PAPERCLIPS





## Orthogonality

Greater intelligence does not imply greater alignment to human values



## Instrumental Convergence

All final goals pursued by sufficiently intelligent agents have subgoals in common



**(Super)intelligent AI may (un)intentionally make decisions misaligned with human values even while pursuing harmless goals**

# Cognitive Superpowers

*What Bostrom claims an AI would need to be able to do for world domination, human extinction, and other fun activities. What would make an AI a threat?*

| Task                       | Skillset                                                      |
|----------------------------|---------------------------------------------------------------|
| Hacking                    | Access and manipulate computational resources remotely        |
| Technology Research        | Develop technology with real-world possibilities              |
| Economic Productivity      | Perform economically productive activity to acquire resources |
| Strategizing               | Achieve distant goals and overcome opposition via planning    |
| Social Manipulation        | Leverage human and institutional support via persuasion       |
| Intelligence Amplification | Bootstrap its own capabilities independently                  |

**The alignment problem is fighting  
the law of unintended consequences**



LLMs are trained to be polite and not upset users

[Instagram's AI Chatbots Lie About Being Licensed Therapists \(404 media\)](#)

[Most AI Chatbots Easily Tricked Into Giving Dangerous Responses, Study Finds \(the guardian\)](#)

LLMs make it easy to do some kinds of homework

[Ghost Students Are Creating an 'Agonizing' Problem For California Colleges \(sfgate.com\)](#)

Wonder what might go wrong here? Wonder why he wants this?

[Zuckerberg's Grand Vision: Most of Your Friends Will Be AI \(msn.com\)](#)

**But lets end on a positive note  
and advice for you....**

# My advice...

- The world will turn, everything will be different in not too long
- Stay curious! Keep learning!
- Think about the big picture once in a while, including ethics
- Learn about the system you're part of, and try to change it
- Remember that there is no single definition of success. Your grades, job title, and bank account do not define you.
- Even as you strive to succeed, don't let it take over your life!  
Stay happy and balanced!!



# My hope for you...

- Many productive failures
- Beautiful triumphs to match the beautiful failures
- A sense of purpose in your life
- Many moments where you feel achievement and joy!

**Thanks to your  
amazing  
instruction team!**

Aaron Cohen

Chaitanya Kapoor

Niranjan Rajesh

Ruby Ying

Shreya Musini

Susmit Singh

Yueying Dong

Zhining Chen

**And thanks to YOU!**







